

GB21802 - Programming Challenges

Week 4 - Graph Problems (Part I)

Claus Aranha

caranha@cs.tsukuba.ac.jp

College of Information Science

2017-05-19,22

Last updated May 17, 2017

Results for the Previous Week

Here are the results for last week:

Week 3: Dynamic Programming I

Deadline: 5/19/2017, 11:59:59 PM (2 days, 11:29 hours from now)

Problems Solved -- 0P:31, 1P:7, 2P:7, 3P:5, 5P:1, 6P:1, 7P:1, 8P:2,

#	Name	Sol/Sub/Total	My Status
1	Wedding shopping	24/29/55	
2	Jill Rides Again	15/15/55	
3	Largest Submatrix	10/11/55	
4	Is Bigger Smarter?	4/4/55	
5	Murcia's Skyline	5/5/55	
6	Trouble of 13-Dots	5/6/55	
7	Exact Change	4/4/55	
8	Unidirectional TSP	3/3/55	

Great Results!

Pre-class Notes (1/2) – ICPC Dates

The dates for the ICPC contest this year are as follows:

- Registration Deadline – 06/30 (Friday)
- National Contest – 07/14 (Friday)

If you want to participate, please talk to me after class or by e-mail. (A team need 3 members)

Pre-class Notes (2/2)

- I have moved the class **Dynamic Programming II** from Week 4 to Week 9;
- The idea is that we will use class 9 to mix different techniques together: (Maths, Graphs, Geometry, DP)
- It will be fun :-)

Week 4 and 5 – Outline

This Week - Graph I

- Graph Basics review: Concepts and Data Structure;
- Depth First Search and Breadth First Search;
- Problems you solve with DFS and BFS;
- Minimum Spanning Tree: Kruskal and Prim Algorithms;

Next Week - Graph II


- Single Source Shortest Path (Dijkstra);
- All Pairs Shortest Path (Floyd Warshall);
- Network Flow and related Problems;
- Bipartite Graph Matching and related Problems;

Many variations in graph problems!

Quick Review of Graph Terms (1)

You probably know all of these. If not, ask questions!



- A Graph G is made of a set of **vertices** V and **edges** E .
- Edges can be **directed** (has source and destination vertices);
- Edges can be **weighted** or not (all weights = 1);
- Sets of nodes can be **connected** or **disconnected**
- Directed Graphs can be **Strongly Connected**
- Edges can be **self-edges**, and/or **multiple edges**



Quick Review of Graph Terms (2)

You probably know all of these. If not, ask questions!



- A **path** is a set of vertices connected by edges;
- A **cycle** is a path with first and last vertices identical;
- **Labelled** graphs and **Isomorphic** graphs;
- A **tree** is a acyclical, undirected graph;
- A **spanning tree** is a subset of edges from E' that form a tree, connecting all nodes $V \in G$;
- A **spamming tree** houses very noisy insects in summer;



Quick Review of Graph Terms (3)

You probably know all of these. If not, ask questions!

- The **degree** of a node is the number of edges connected to it;
- Directed graphs have **in-degrees** and **out-degrees**;
- A **bipartite** graph can be divided in two sets of unconnected vertices;
- A **Match** or **Pairing** is a set of edges that connects the nodes in the bipartite graph;



Data Structures for Graphs (1)

Adjacency Matrix - Stores connection between Vertices

```
int adj[100][100];  
// adj[i][j] is 0 if no edge between i,j  
// adj[i][j] is A if edge of weight A links i,j
```

- **Pro:** Very simple to program, manipulate;
- **Con:** Cannot store multigraph; Wastes space for sparse graphs;
Requires time $O(V)$ to calculate number of neighbors;

Edge List – Stores Edges list for each Vertex

```
typedef pair<int,int> ii;  
typedef vector<ii> vii;  
vector<vii> AdjList;
```

- **Pro:** $O(V + E)$ space, efficient if graph is sparse; Can store multigraph;
- **Con:** A (bit) more code than Adjacency Matrix

Data Structures for Graphs (2)

Edge List

```
vector< pair <int,ii>> Edgelist;
```

Stores a list of all the edges in the graph. Vertices are implicit from the edge list. This is useful for Kruskal's algorithm (which we will see later), but otherwise complicates things.

Implicit Graph

Some graphs **do not** need to be stored in a special structure if they have very clear rules about when two vertices connect.

Examples:

- A square grid;
- Knight's chess moves;
- Two vertices i, j connect if $i + j$ is prime;

Searching in a Graph: BFS and DFS

Almost all graph problems involve visiting each of its vertices in some form. There are two approaches for visiting the nodes in a graph:

Depth First Search – DFS


DFS is commonly implemented as a recursive search. For every node visited, immediately visit the first edge in it, backtracking when a loop is reached, or no more edges can be followed.

Breadth First Search – BFS


BFS is commonly implemented iterating over a FIFO queue. For every node visited, all new edges are put on the back of the queue. Visit the next edge at the top of the queue.

BFS/DFS: Visualization

DFS




BFS




BFS/DFS: Visualization



DFS



BFS




BFS/DFS: Visualization


DFS**BFS**

BFS/DFS: Visualization

DFS




BFS




BFS/DFS: Visualization

DFS




BFS




BFS/DFS: Visualization

DFS




BFS




BFS/DFS: Visualization

DFS




BFS




BFS/DFS: Visualization

DFS



BFS



BFS/DFS: Implementation

There are many ways to implement BFS/DFS, here is a suggestion.

DFS

```
vector<int> dfs_vis; // initially all set to UNVISITED
void dfs(int v) {
    dfs_vis[v] = VISITED;
    for (int i; i < (int)Adj_list[v].size(); i++) {
        pair <int,int> u = Adj_list[v][i];
        if (dfs_vis[u.first] == UNVISITED) dfs(u.first)
    }
}
```


BFS

```
vector<int> d(V,INF); d[s] = 0; queue<int> q; q.push(s);
while(!q.empty()) {
    u = q.front(); q.pop();
    for (int i=0; i < (int)Adj_list[u].size(); i++) {
        pair <int,int> v = Adj_list[u][i]; //same as dfs
        if (d[v.first] == INF) {
            d[v.first] = d[u] + 1; q.push(v.first);
        }
    }
}
```

Simple BFS/DFS – UVA 11902: Dominator

Problem Summary

Vertex X **dominates** vertex Y if every path from a start vertex 0 to Y must go through X . Determine which nodes dominate which other.




- 0 dominates all nodes;
- 3 dominates 4;
- 1 does not dominate 3;

How do you solve it?

Simple BFS/DFS – UVA 11902: Dominator

Solution

```
DFS(0);  
for i in (0:N):  
    if i is reached: dominate[0][i] = 1;  
  
for i in (1:N):  
    remove i from graph;  
    DFS(0)  
    for j in (1:N):  
        if (j is not reached) and (dominate[0][j] == 1):  
            dominate[i][j] = 1  
    return i to graph
```




Common Algorithms: Connected Components

With small modifications to BFS/DFS, we can solve many simple problems

Since a single run of DFS/BFS finds all connected nodes, we can use it to find (and count) all the connected components (CC) of an **undirected** graph.

```
numCC = 0;
dfs_num.assign(V, UNVISITED);
for (int i = 0; i < V; i++)
    if (dfs_num[i] == UNVISITED)
        cout << "\nCC " << ++numCC << ":"; dfs(i);
// modify dfs() to print every node it visits
```



CC 1: 0 1 2

CC 2: 3 5

CC 3: 4

Flood Fill

A simple tweak of the BFS (or DFS) can be used to **label/color** and count the size of each CC.

“flood fill” is often used in problems involving implicit 2D grids.

```
####..#  
#.###.#  
#..@.##  
##d.###  
#..####
```

```
int dr[] = {1,1,0,-1,-1,-1,0,1}; // trick to explore an  
int dc[] = {0,1,1,1,0,-1,-1,-1}; // implicit NESW graph  
  
int floodfill(int y, int x, char c1, char c2) {  
    if (y < 0 || y >= R || x < 0 || x >= C) return 0;  
    if (grid[y][x] != c1) return 0;  
    int ans = 1;  
    grid[y][x] = c2;  
    for (int d = 0; d < 8; d++)  
        ans += floodfill(y+dr[d], x+dc[d], c1, c2);  
    return ans;  
}
```


Topological Sort (Directed Acyclic Graphs)

A Topological sort is a linear ordering of vertices of a DAG so that vertex u comes before vertex v if edge $u \rightarrow v$ exists in the DAG. Topological Sorts are useful for problems involving the ordering of pre-requisites.

Khan's algorithm for Topological sort (modified edge-BFS)

```
Q = queue(); toposort = list();
for j in edge:
    in_degree[j.destination] += 1
for i in node:
    if in_degree[i] == 0: Q.add(i);
while (Q.size() > 0):
    u = Q.dequeue(); toposort.add(u);
    for i in u.out_edges():
        v = i.destination
        in_degree[v] -= 1
        if in_degree[v] == 0:
            Q.add(v);
```


Bipartite Check


To check whether a graph is bipartite, we perform a BFS or DFS on the graph, and set the color of every node to black or white, alternatively. Pay attention to collision conditions.

```
queue<int> q; q.push(s);
vector<int> color(V, INF); color[s] = 0;
bool isBipartite = true;
while (!q.empty() && isBipartite) {
    int u = q.front(); q.pop();
    for (int j=0; j < adj_list[u].size(); j++) {
        pair<int,int> v = adj_list[u][j];
        if (color[v] == INF) {
            color[v.first] = 1 - color[u];
            q.push(v.first);
        }
        else if (color[v.first] == color[u]) {
            isBipartite = False;
        }
    }
}
```

Bipartite Check – Visualization


Testing Bipartite property






Bipartite Check – Visualization

Testing Bipartite property




Bipartite Check – Visualization

Testing Bipartite property




Bipartite Check – Visualization

Testing Bipartite property




Bipartite Check – Visualization

Testing Bipartite property



Rearranging the nodes




Articulation Points and Bridges

Problem Description

In an undirected graph G :

- A vertex V is an **Articulation Point** if removing V would make G disconnected.
- An edge E is a **Bridge** if removing E would make G disconnected.



Articulation Points and Bridges: Algorithm

Complete Search algorithm for Articulation Points

- 1 Run DFS/BFS, and count the number of CC in the graph;
- 2 For each vertex v , remove v and run DFS/BFS again;
- 3 If the number of CC increases, v is a connection point;


Since DFS/BFS is $O(V + E)$, this algorithm runs in $O(V^2 + EV)$.

... but we can do better!

Tarjan's DFS variant for Articulation point ($O(V+E)$)


Tarjan Variant: $O(V + E)$

Main idea: Add extra data to the DFS to detect articulations.



- $dfs_num[]$: Receives the number of the iteration when this node was reached for the first time;
- $dfs_low[]$: Receives the lowest $dfs_num[]$ which can be reached if we start the DFS from here;
- For any neighbors u, v , if $dfs_low[v] \geq dfs_num[u]$, then u is an articulation node.

Tarjan's DFS variant for Articulation point (2)



```
void dfs_a(u) {
    dfs_num[u] = dfs_low[u] = IterationCounter++; // dfs_num[u] is a simple counter
    for (int i = 0; i < AdjList[u].size(); i++) {
        v = AdjList[u][i];
        if (dfs_num[v] == UNVISITED) {
            dfs_parent[v] = u; // store parent
            if (u == 0) rootChildren++; // special case for root node

            dfs_a(v);
            if (dfs_low[v] >= dfs_num[u])
                articulation_vertex[u] = true;
            dfs_low[u] = min(dfs_low[u], dfs_low[v])


        } else if (v != dfs_parent[u]) // found a cycle edge
            dfs_low[u] = min(dfs_low[u], dfs_num[v])
    }
}
```

Strongly Connected Components (Directed Graph)


Problem Description

On a directed graph G , a Strongly Connected Component (SCC) is a subset G' where for every pair of nodes $a, b \in G'$, there is both a path $a \rightarrow b$ and a path $b \rightarrow a$.

One CC



Three SCC



Strongly Connected Components – Algorithm

We can use a simple modification of the algorithm for bridges and articulation points:




- Every time we visit a new node, put that node in a stack S ;
- When we finish visiting a node i , test if $\text{dfs_num}[i] == \text{dfs_min}[i]$.
- If the above condition is true, i is the root of the SCC. Pop all vertices in the stack as part of the SCC.

Minimum Spanning Trees (MST)

Definition

A **Spanning Tree** is a subset E' from graph G so that all vertices are connected without cycles.

A **Minimum Spanning Tree** is a spanning tree where the sum of edge's weights is minimal.



MST – Use cases and Algorithms

Problems using MST

Problems using MST usually involve calculating the minimum costs of infrastructure such as roads or networks.

Some variations may require you to find the **maximum** spanning tree, or define some edges that **must** be taken in advance.

Algorithms for MST

The two main algorithms for calculating the MST are the **Kruskal's** algorithms and the **Prim's** algorithms.


Both are greedy algorithms that add edges to the MST in weight order.

Kruskal's Algorithm

Outline

Kruskal's algorithm sorts all edges by their weight, and try to add each edge to the MST, checking whether adding that edge would create a cycle.

- 1 Sort all edges;
- 2 If smallest edge does not create a cycle, add to MST;
- 3 If smallest edge creates a cycle, remove it from list;
- 4 Go to 2;




Kruskal's Algorithm

Outline

Kruskal's algorithm sorts all edges by their weight, and try to add each edge to the MST, checking whether adding that edge would create a cycle.

- 1 Sort all edges;
- 2 If smallest edge does not create a cycle, add to MST;
- 3 If smallest edge creates a cycle, remove it from list;
- 4 Go to 2;




Kruskal's Algorithm

Outline

Kruskal's algorithm sorts all edges by their weight, and try to add each edge to the MST, checking whether adding that edge would create a cycle.

- 1 Sort all edges;
- 2 If smallest edge does not create a cycle, add to MST;
- 3 If smallest edge creates a cycle, remove it from list;
- 4 Go to 2;




Kruskal's Algorithm

Outline

Kruskal's algorithm sorts all edges by their weight, and try to add each edge to the MST, checking whether adding that edge would create a cycle.

- 1 Sort all edges;
- 2 If smallest edge does not create a cycle, add to MST;
- 3 If smallest edge creates a cycle, remove it from list;
- 4 Go to 2;




Kruskal's Algorithm

Outline

Kruskal's algorithm sorts all edges by their weight, and try to add each edge to the MST, checking whether adding that edge would create a cycle.

- 1 Sort all edges;
- 2 If smallest edge does not create a cycle, add to MST;
- 3 If smallest edge creates a cycle, remove it from list;
- 4 Go to 2;



Kruskal's Algorithm – Implementation

```
vector<pair<int, pair<int,int>> Edgelist;
sort(Edgelist.begin(),Edgelist.end());
int mst_cost = 0;
UnionFind UF(V);
    // note 1: Pair object has built-in comparison;
    // note 2: Need to implement UnionSet class;

for (int i = 0; i < Edgelist.size(); i++) {
    pair <int, pair <int,int>> front = Edgelist[i];
    if (!UF.isSameSet(front.second.first,
                      front.second.second)) {
        mst_cost += front.first;
        UF.unionSet(front.second.first,front.second.second)
    }
}


cout << "MST Cost: " << mst_cost << "\n"
```

Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;




Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;




Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;




Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;




Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;




Prim's Algorithm

Outline

Prim's algorithm adds nodes to the MST one at a time, and keeps the edges connected to those nodes in a **priority queue**. It then tests each edge in the priority queue to add more nodes to the MST, avoiding cycles.

- 1 Add node 0 to MST;
- 2 Add all edges from new node to Priority Queue;
- 3 Visit smallest edge in Queue;
- 4 If the edge leads to a new node, add it to MST;
- 5 Add new edges to Queue;
- 6 Go to 3;



Prim's Algorithm – Implementation

```
vector <int> taken;
priority_queue <pair <int,int>> pq;


void process (int v) {
    taken[v] = 1;
    for (int j = 0; j < (int)AdjList[v].size(); j++) {
        pair <int,int> ve = AdjList[v][j];
        if (!taken[ve.first])
            pq.push(pair <int,int> (-ve.second,-ve.second))
    }
    taken.assign(V,0);
    process(0);
    mst_cost = 0;

    while (!pq.empty()) {
        vector <int,int> pq.top(); pq.pop();
        u = -front.secont, w = -front.first;
        if (!taken[u]) mst_cost += w, process(u);
    }
}
```

MST variant 1 – Maximum Spanning tree

The **Maximum Spanning Tree** variant requires the spanning tree to have maximum possible weight.


It is very easy to implement the Maximum MST by reversing the sort order of the edges (Kruskal), or the weighting of the priority Queue (Prim).



MST variant 2 – Minimum Spanning Subgraph, Forest

In one importante variant of the MST, a subset of edges or vertices are pre-selected.

- In the case of pre-selected vertices, add them to the “taken” list in Kruskal’s algorithm before starting;
- In the case of edges, add the end vertices to the “taken” list;
- What if you are given a **number of Connected Components**?



MST Variant 3 – n th Best MST


Problem Definition

Consider that you can order MST by their costs: G_1, G_2, \dots, G_n . This variant asks you to calculate the n^{th} best spanning tree.

Basic Idea:

- Calculate the MST (using Kruskal or Prim);
- For every edge in the MST, remove that edge from the graph and calculate a new MST;
- The new MST with minimum weight is the 2^{nd} best MST;

MST Variant 4 – Min-max (or Max-min)



Problem Definition

Given two vertices i, j , find a path $i \rightarrow j$ so that the cost of the most expensive edge is minimized;

Another way to write this problem is: Find the cheapest path where the cost of the path is the cost of the most expensive edge.

How to solve

The MST finds the path that connects all nodes while keeping the cost of individual edges minimal.

To solve the minimax problem, we calculate the MST for G , and then find the path from i to j in the MST.

Summary

- Graphs come in a wide variety of types;
- Graph problems also have many different types;
- Most problems involve small modifications of DFS and BFS;

Next Week

More Graphs!

- Shortest Paths (Single Source and All Pairs);
- Network Flow (and related problems);
- Graph Matching (bipartite matching, etc) (and related problems);

This Week's Problems

- Dominator
- Forwarding Emails
- Ordering
- Place the Guards
- Doves and Bombs
- Come and Go
- ACM Contest and Blackout
- Ancient Messages

Problem Hints (0)

Library!

For many of these problems, you will use a lot of repeated code:

- Visited Node arrays;
- Adjacent lists;
- Parent nodes;

Prepare a template for the most common codes you use, and copy-paste it whenever necessary!

Tricky Cases

- Graphs with 1 or 0 Vertex
- Unconnected Graphs
- Self Loops
- Double edges

Problem Hints (1)

Dominator

- If All paths from 0 to node B pass through node A, then node A **dominates** node B;
- For all pair of nodes i, j , output “Y” if i **dominates** j , or “N” if not;

The idea of this problem is one of “reachability” – can I reach node j if I remove node i from the graph?

Note: if j is not connected to “0”, then *no one dominates j*

Problem Hints (2)

Forwarding Emails

Every person i sends e-mail only to person j .

What is the longest email chain?

Where does it start?

- How do you deal with loops?
- Time limit is not very large, Try to find an $O(n)$ solution!

Problem Hints (3)

Ordering

Print all possible Orderings of a Direct Acyclic Graph

Generalize the DAG ordering algorithm which we discussed in class.

Palace Guards

- How do you represent the roads and junctions as a Graph?
- Find a “guard-no guard” assignment to vertices of the graph.
- First test if a solution is possible!

Problem Hints (4)

Doves and Bombs

This problem is about finding “critical vertices” in a graph. But how do you calculate the “pigeon value” of a vertex?

Come and Go

Straight implementation of “Strong Connected Components”. Be careful with tricky graphs!

Problem Hints (5)

ACM Contest and Blackout

Goal: Find the **First** minimum spanning Tree and the **Second** minimum spanning Tree

- In this class we discussed how to find the Minimum Spanning Tree
- How would we find the **second minimum**?
- Idea: Maybe if we remove some edges from the graph?

Problem Hints (6)

Ancient Message – Challenge problem!

Count the symbols inside an image – order does not matter!

What is the **Main** difference between the symbols?

- The shape and size of the symbols is actually not important!
- Before you begin programming, discover what is the real difference between the symbols.
- Hint: The numbers “1”, “0”, “8” have the same difference.